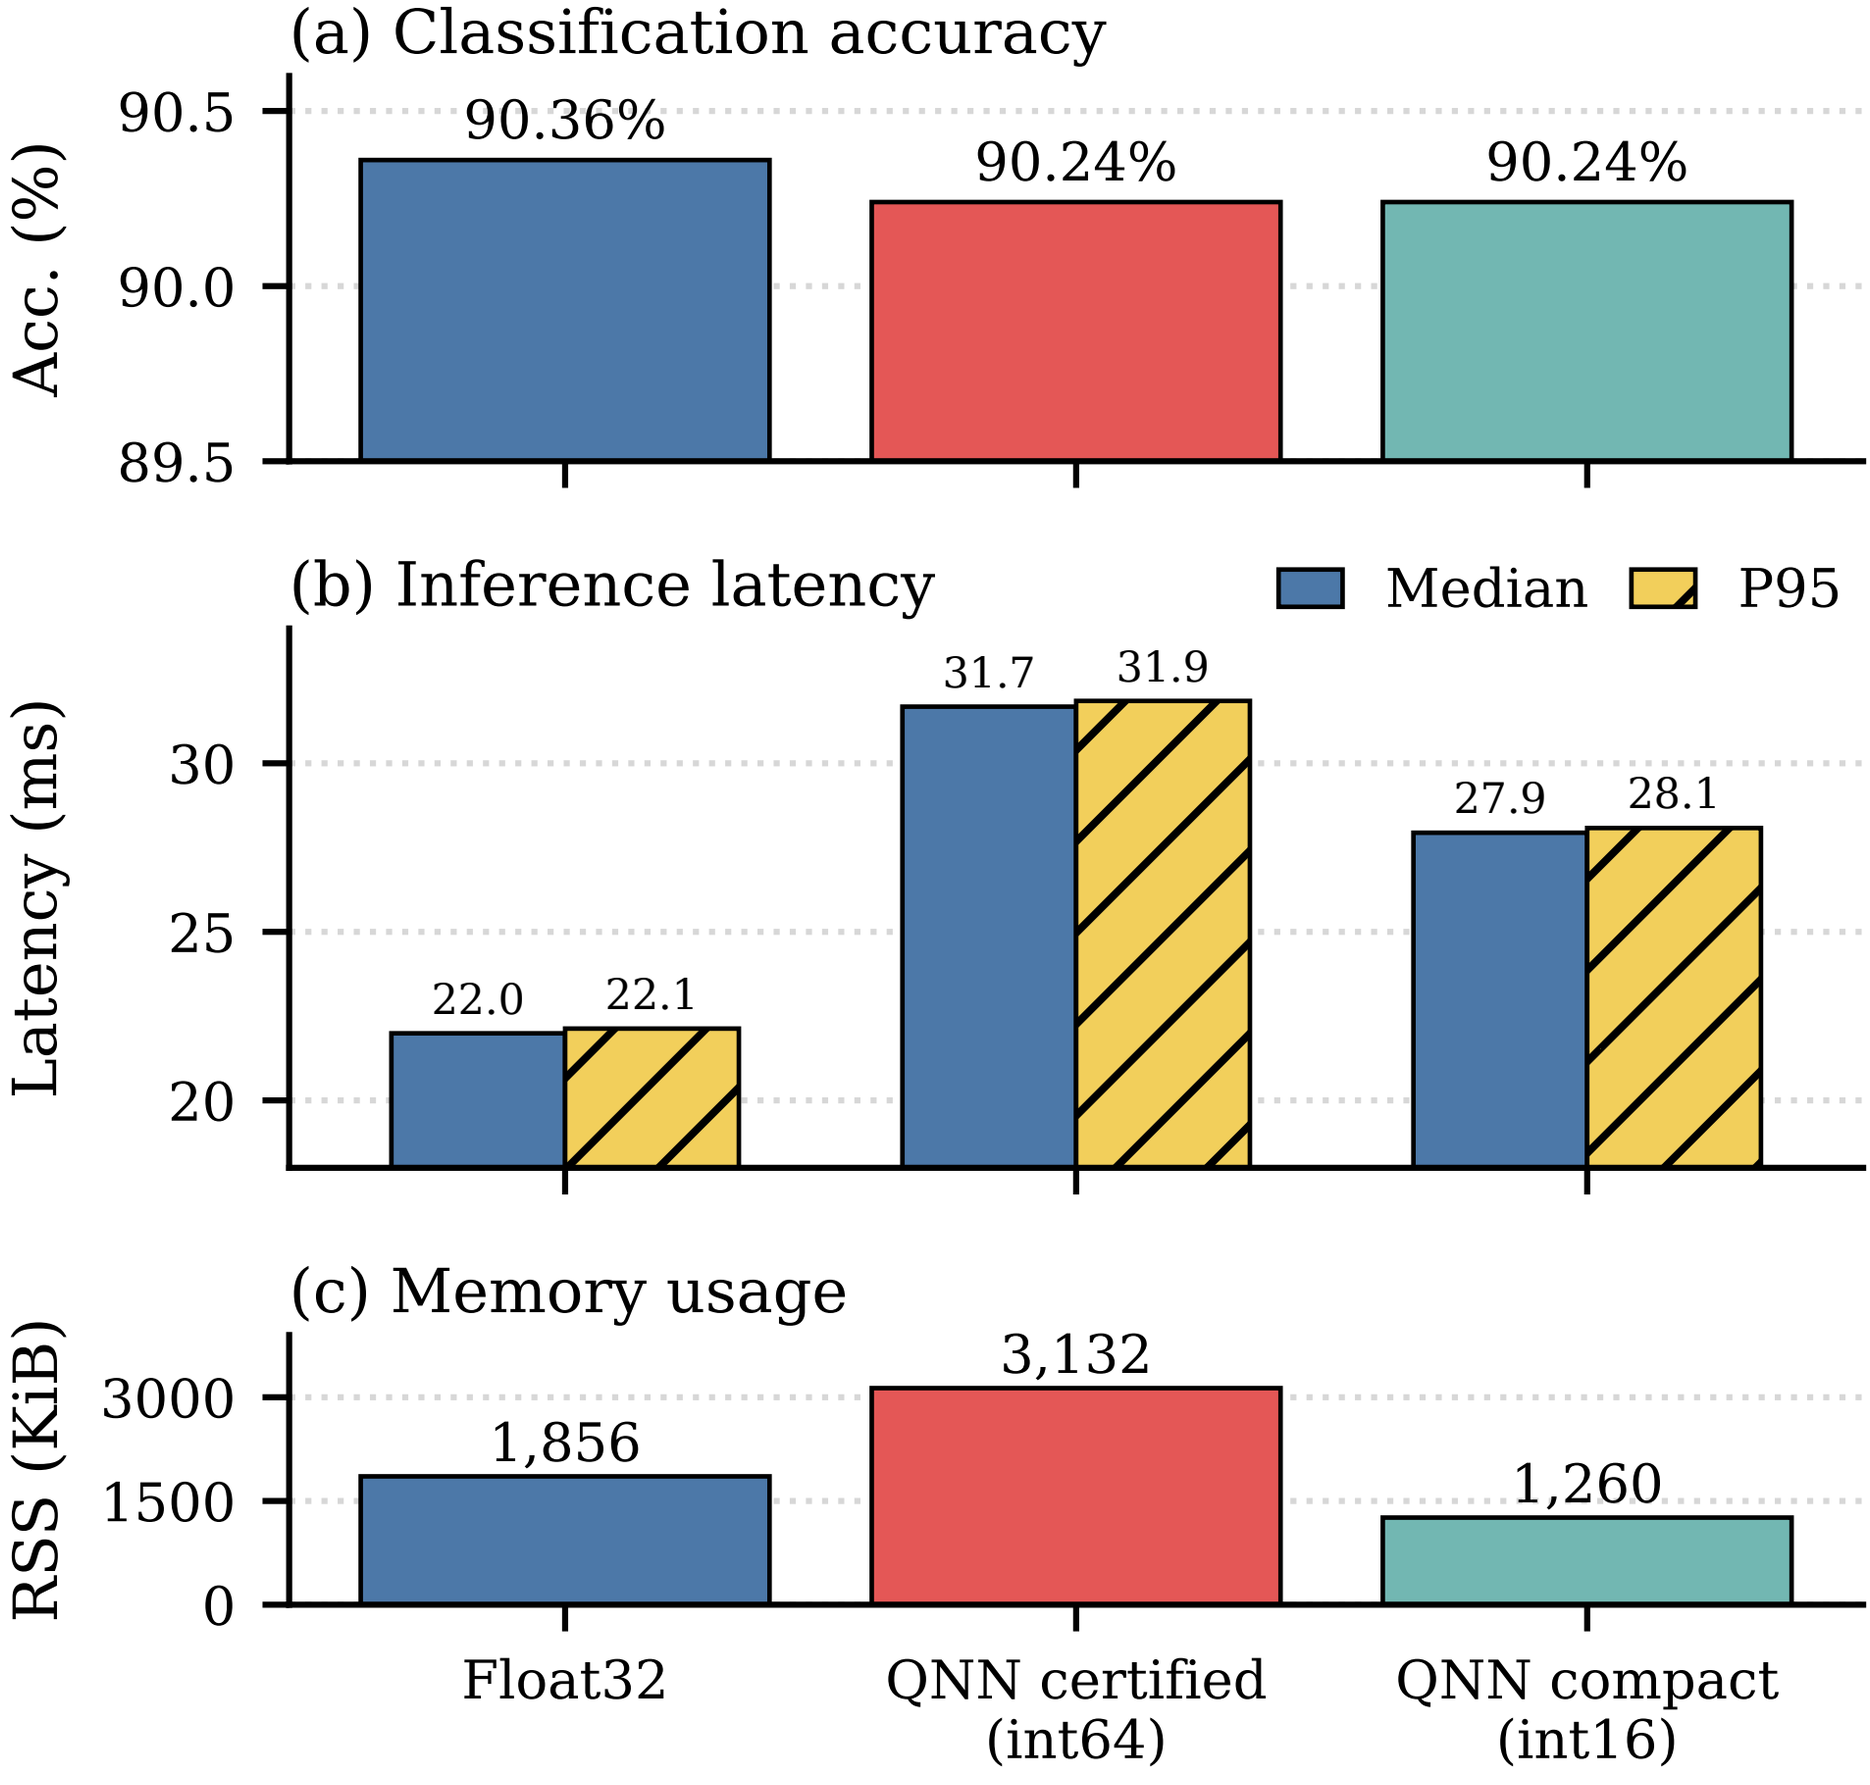
(a) Classification accuracy
Acc. (%)
90.5
90.0
89.5
90.36%
90.24%
90.24%
(b) Inference latency
Median
P95
Latency (ms)
30
25
20
31.7
31.9
27.9
28.1
22.0
22.1
(c) Memory usage
RSS (KiB)
3000
1500
0
1,856
3,132
1,260
Float32
QNN certified (int64)
QNN compact (int16)